

# An introduction to modern tools for collaborative science

## Lecture 4 - Docker and GitHub actions

Luca Heltai <[luca.heltai@sissa.it](mailto:luca.heltai@sissa.it)>

Mathematical Analysis, Modeling, and Applications ([math.sissa.it](http://math.sissa.it))

Theoretical and Scientific Data Science ([datascience.sissa.it](http://datascience.sissa.it))



# docker

## Introduction to Docker

Adapted from official [docker.io](https://docker.io) slides of November, 2013

Luca Heltai <[luca.heltai@sissa.it](mailto:luca.heltai@sissa.it)>

# The Challenge

Multiplicity of Stacks



Static website

nginx 1.5 + modsecurity + openssl + bootstrap 2



Background workers

Python 3.0 + celery + pyredis + libcurl + ffmpeg + libopencv + nodejs + phantomjs



User DB

postgresql + pgv8 + v8



Queue

Redis + redis-sentinel



Analytics DB

hadoop + hive + thrift + OpenJDK



Web frontend

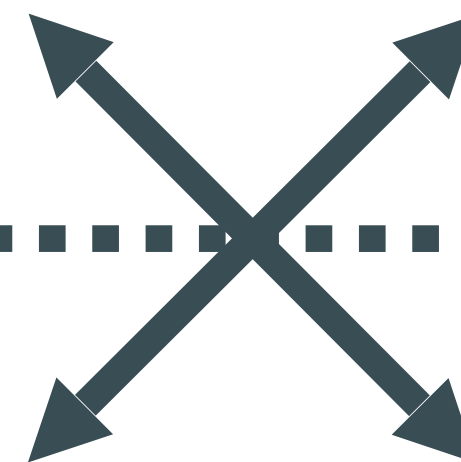
Ruby + Rails + sass + Unicorn



API endpoint

Python 2.7 + Flask + pyredis + celery + pycopg + postgresql-client

Do services and apps  
interact  
appropriately?



Multiplicity of  
hardware  
environments



Development VM

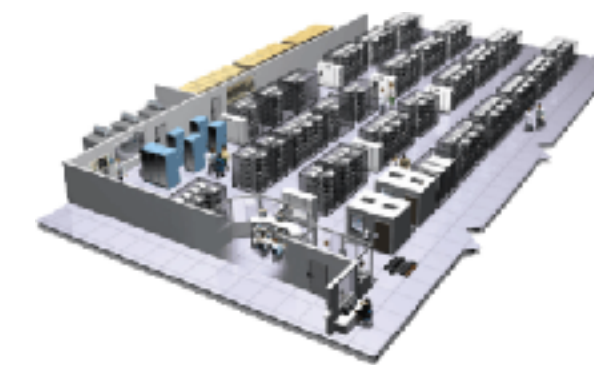


QA server

Customer Data Center



Public Cloud



Disaster recovery

Production Servers

Production Cluster









Contributor's laptop



Can I migrate  
smoothly and  
quickly?



# The Matrix From Hell

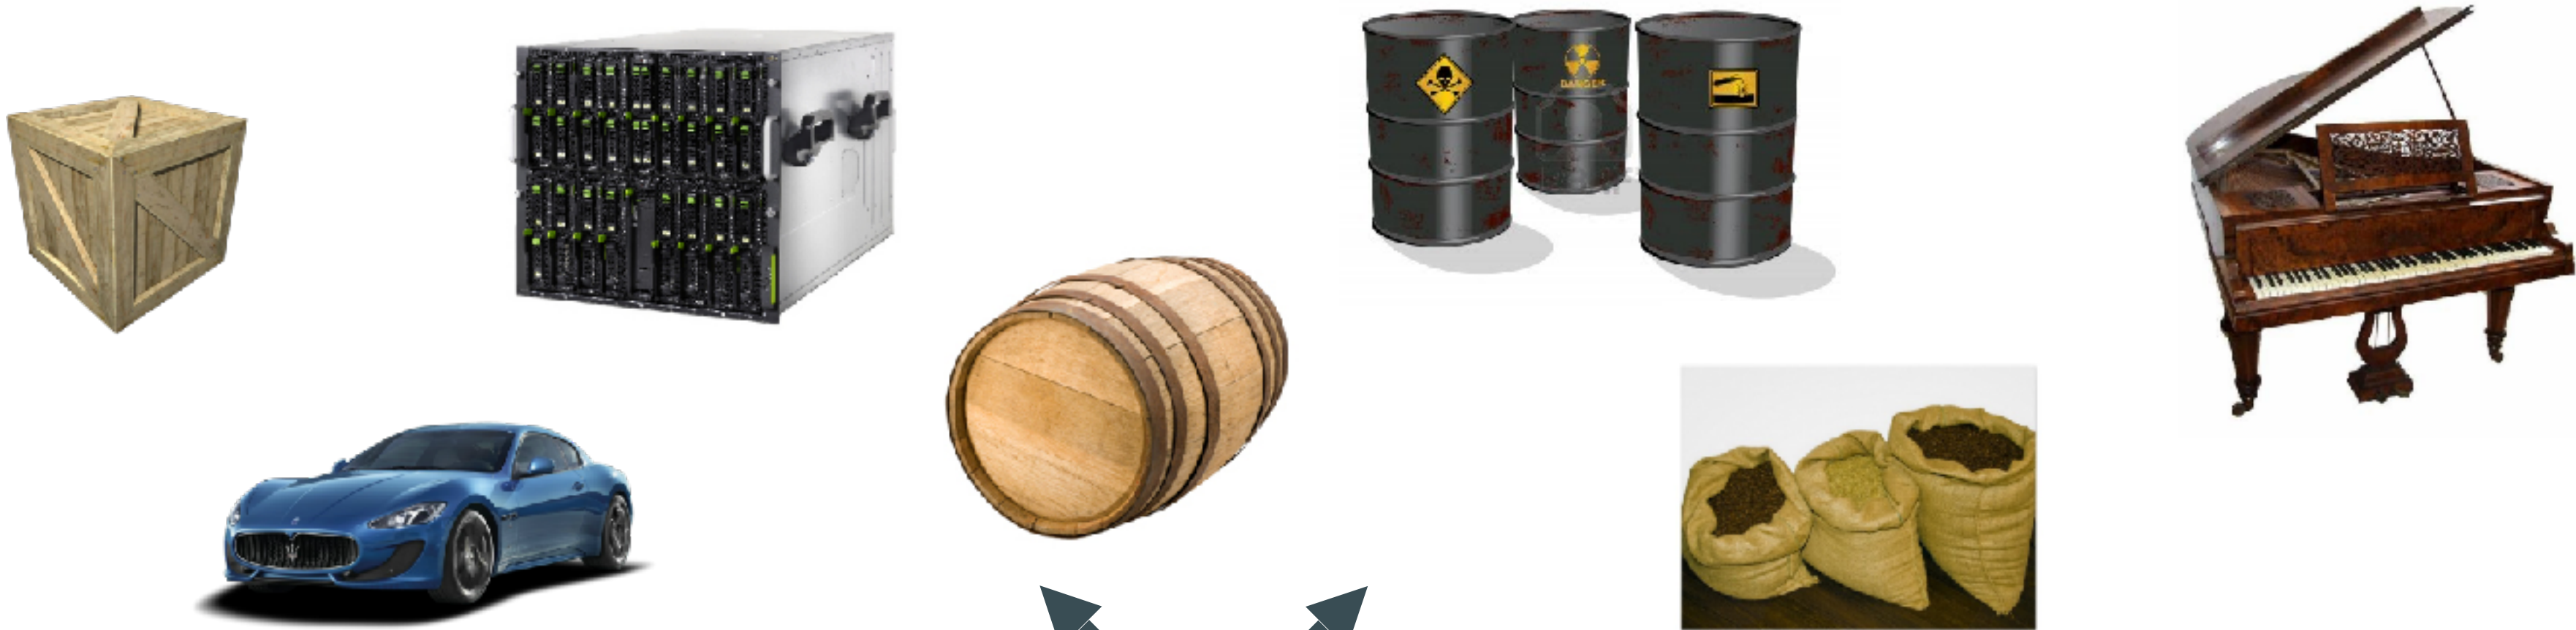
	Static website	?	?	?	?	?	?	?
	Web frontend	?	?	?	?	?	?	?
	Background workers	?	?	?	?	?	?	?
	User DB	?	?	?	?	?	?	?
	Analytics DB	?	?	?	?	?	?	?
	Queue	?	?	?	?	?	?	?
		Development VM	QA Server	Single Prod Server	Onsite Cluster	Public Cloud	Contributor's laptop	Customer Servers





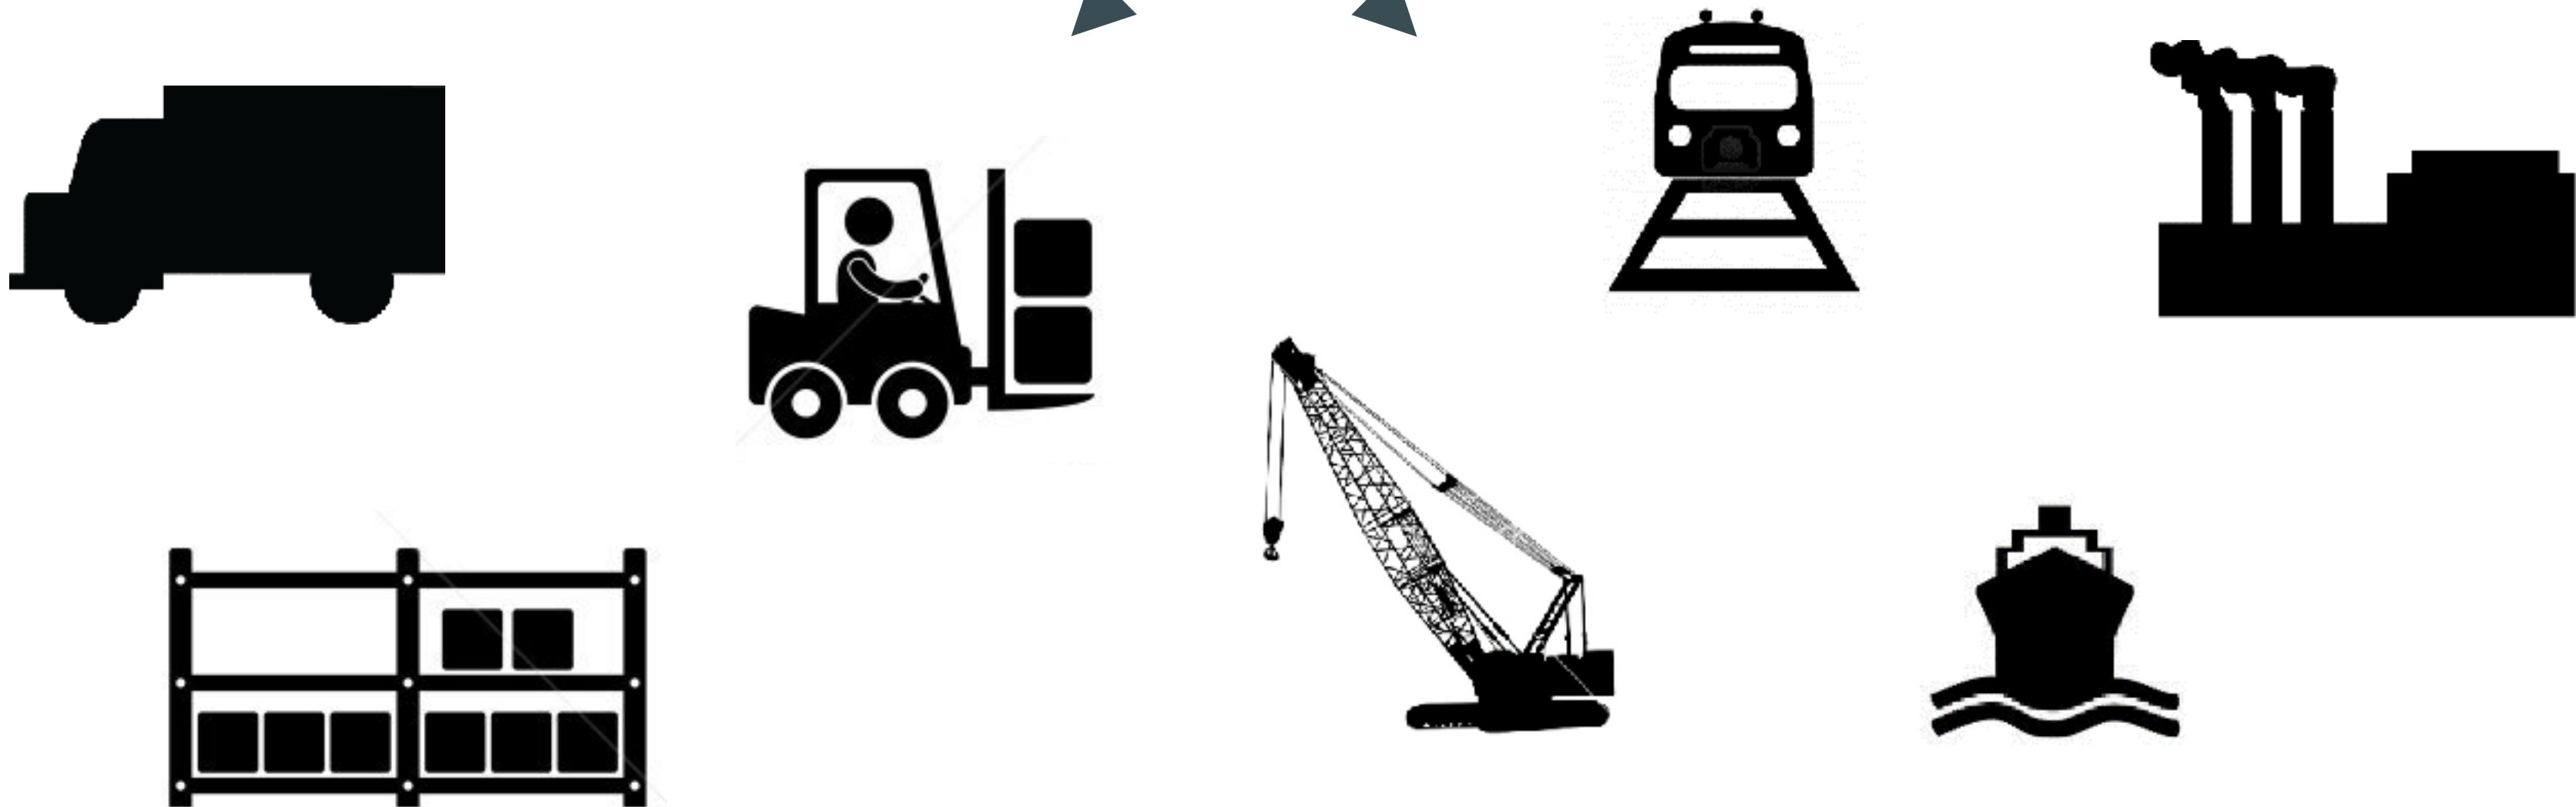
# Cargo Transport Pre-1960

Multiplicity of Goods










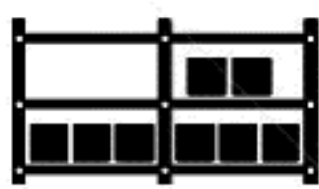





Do I worry about how goods interact (e.g. coffee beans next to spices)

Multiplicity of methods for transporting/storing



Can I transport quickly and smoothly (e.g. from boat to train to truck)

# Also a matrix from hell

	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
							



# Solution: Intermodal Shipping Container

Multiplicity of Goods



A standard container that is loaded with virtually any goods, and stays sealed until it reaches final delivery.

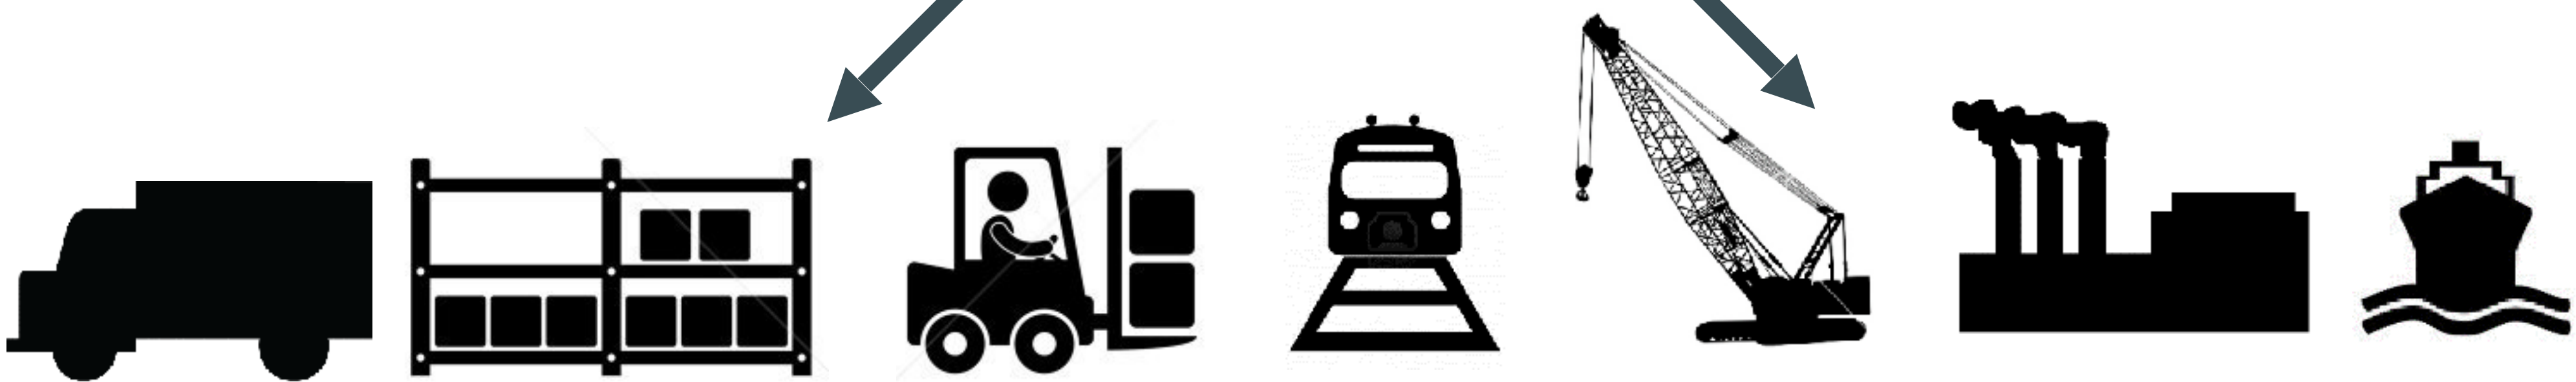
Do I worry about how goods interact (e.g. coffee beans next to spices)



...in between, can be loaded and unloaded, stacked, transported efficiently over long distances, and transferred from one mode of transport to another

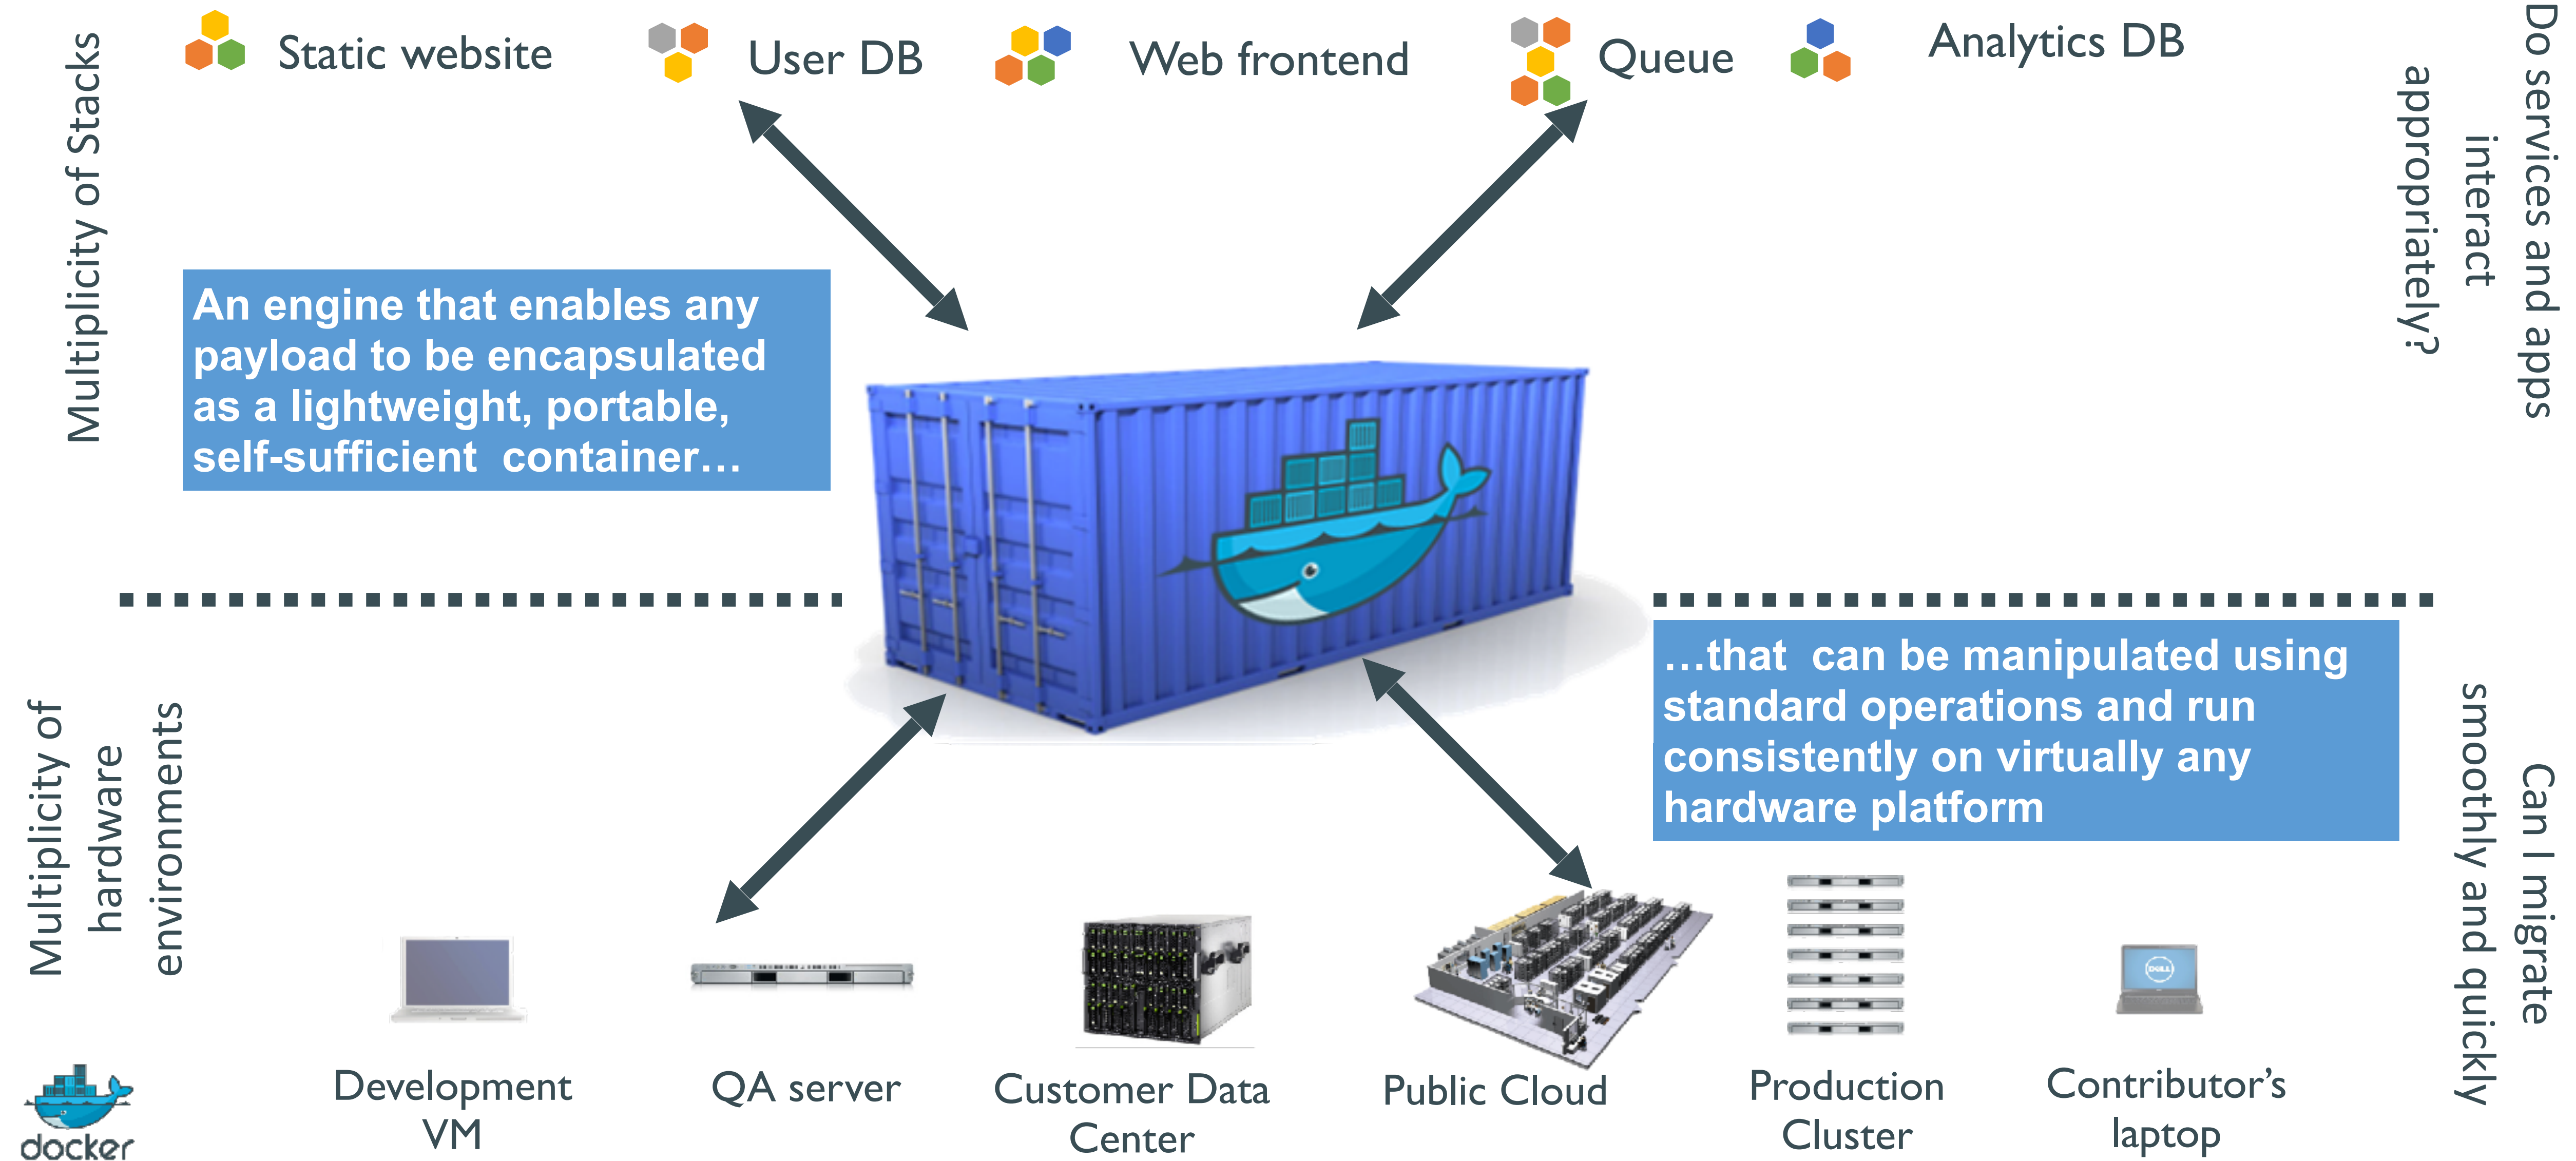
Can I transport quickly and smoothly (e.g. from boat to train to truck)

Multiplicity of methods for transporting/storing





























































# Docker is a shipping container system for code





# Docker eliminates the matrix from Hell

	Static website							
	Web frontend							
	Background workers							
	User DB							
	Analytics DB							
	Queue							
		Development VM	QA Server	Single Prod Server	Onsite Cluster	Public Cloud	Contributor's laptop	Customer Servers
								

# Why Developers Care

---

- Build once...(finally) run anywhere\*
  - A clean, safe, hygienic and portable runtime environment for your app.
  - No worries about missing dependencies, packages and other pain points during subsequent deployments.
  - Run each app in its own isolated container, so you can run various versions of libraries and other dependencies for each app without worrying
  - Automate testing, integration, packaging...anything you can script
  - Reduce/eliminate concerns about compatibility on different platforms, either your own or your customers.
  - Cheap, zero-penalty containers to deploy services? A VM without the overhead of a VM? Instant replay and reset of image snapshots? That's the power of Docker

\* With the 0.7 release, we support any x86 server running a modern Linux kernel (3.2+ generally. 2.6.32+ for RHEL 6.5+, Fedora, & related)



# Why Devops Cares?

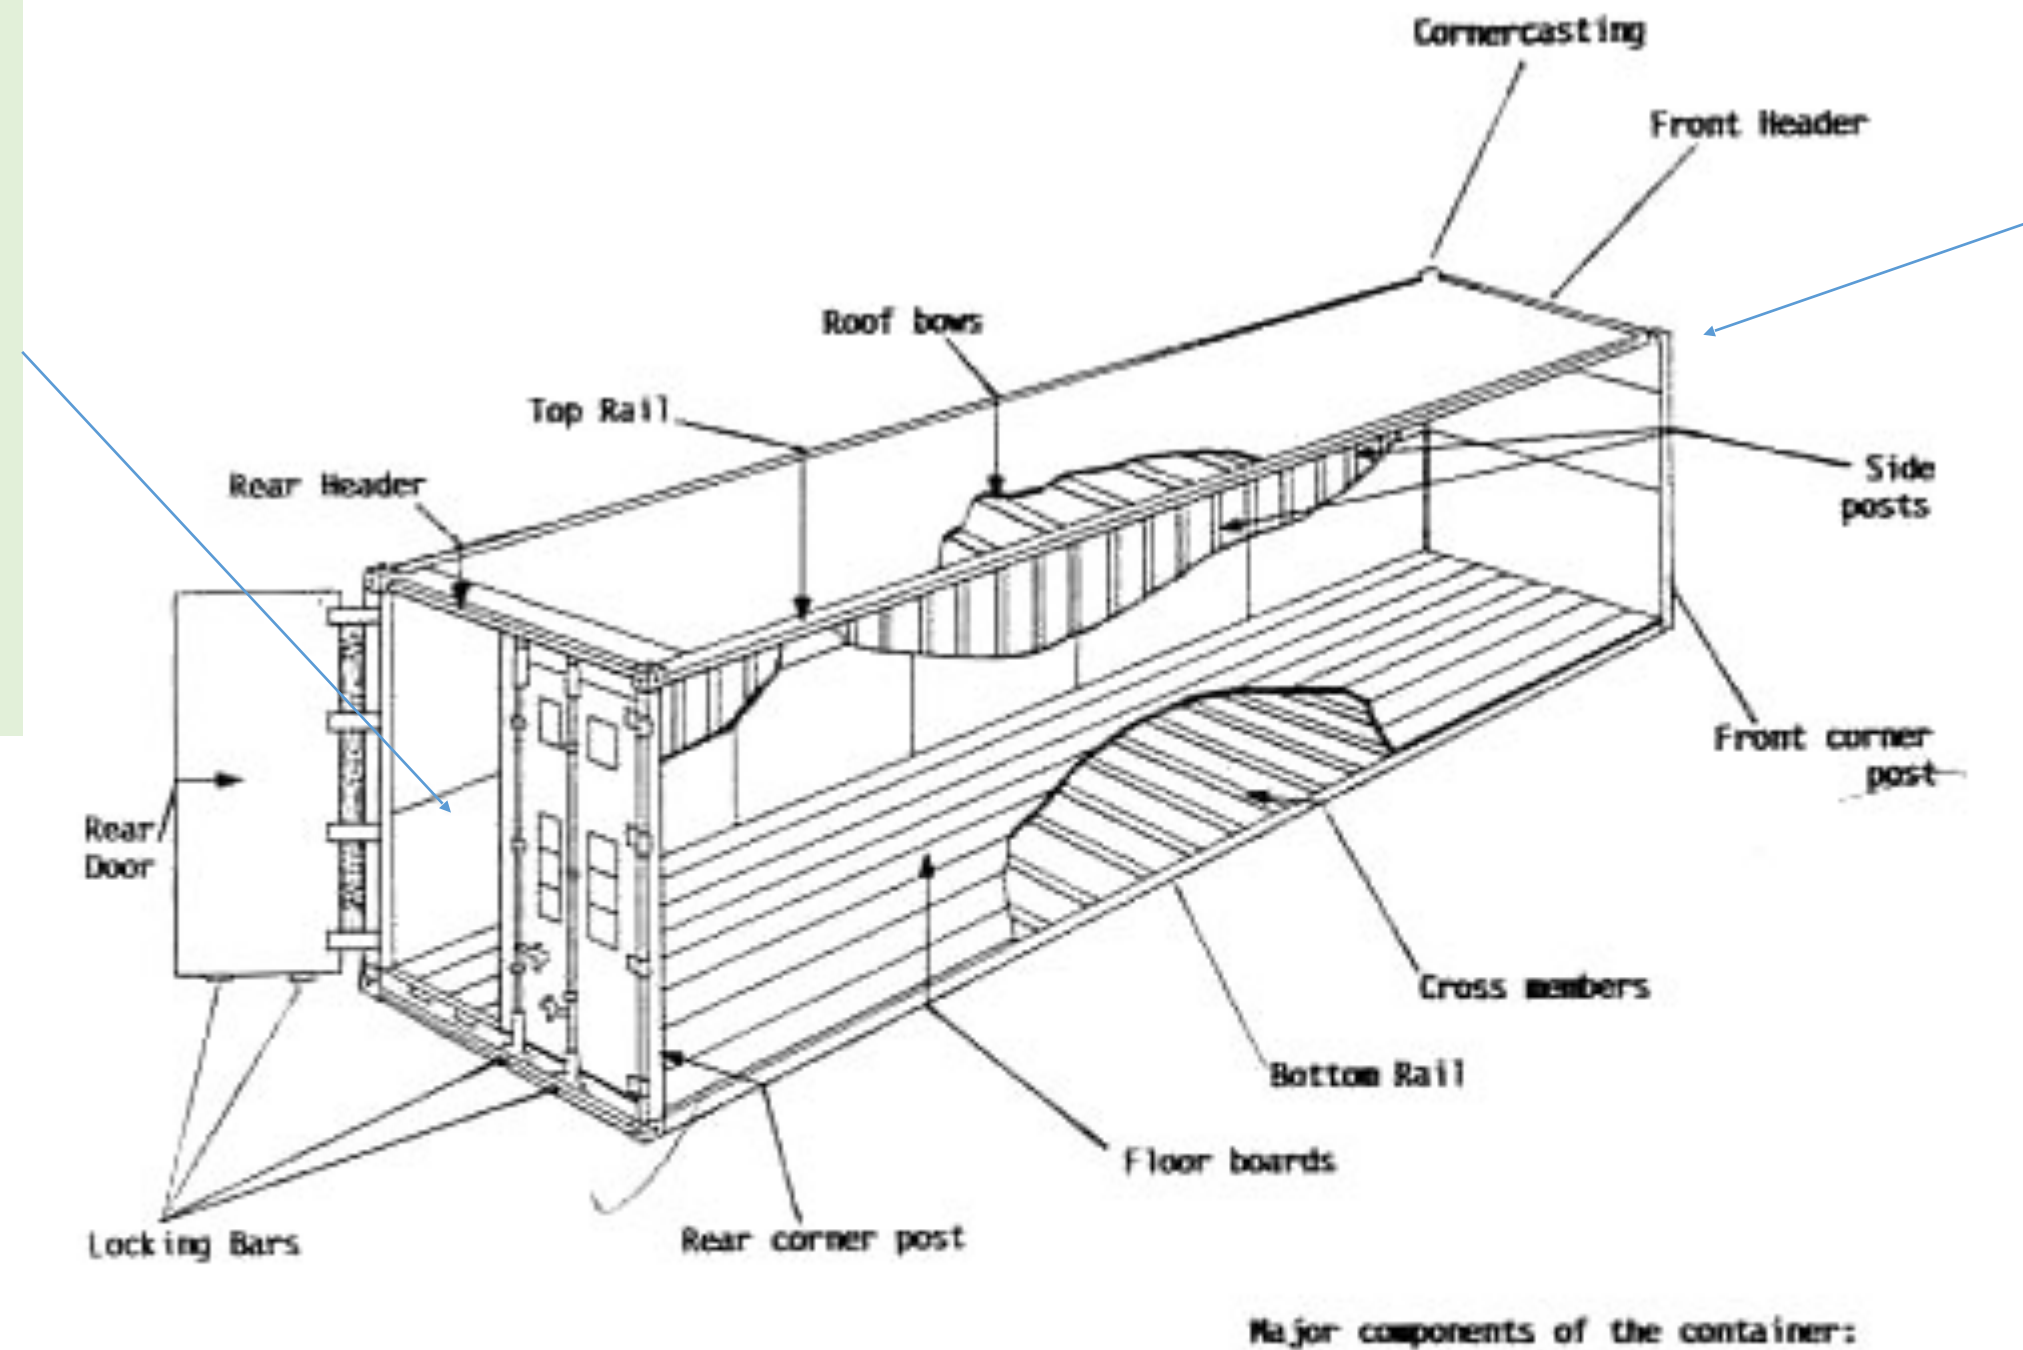
---

- **Configure once...run anything**
  - Make the entire lifecycle more efficient, consistent, and repeatable
  - Increase the quality of code produced by developers.
  - Eliminate inconsistencies between development, test, production, and customer environments
  - Support segregation of duties
  - Significantly improves the speed and reliability of continuous deployment and continuous integration systems
  - Because the containers are so lightweight, address significant performance, costs, deployment, and portability issues normally associated with VMs



# Why it works—separation of concerns

- Dan the Developer
  - Worries about what's "inside" the container
    - His code
    - His Libraries
    - His Package Manager
    - His Apps
    - His Data
  - All Linux servers look the same



- Oscar the Ops Guy
  - Worries about what's "outside" the container
    - Logging
    - Remote access
    - Monitoring
    - Network config
  - All containers start, stop, copy, attach, migrate, etc. the same way

# More technical explanation

---

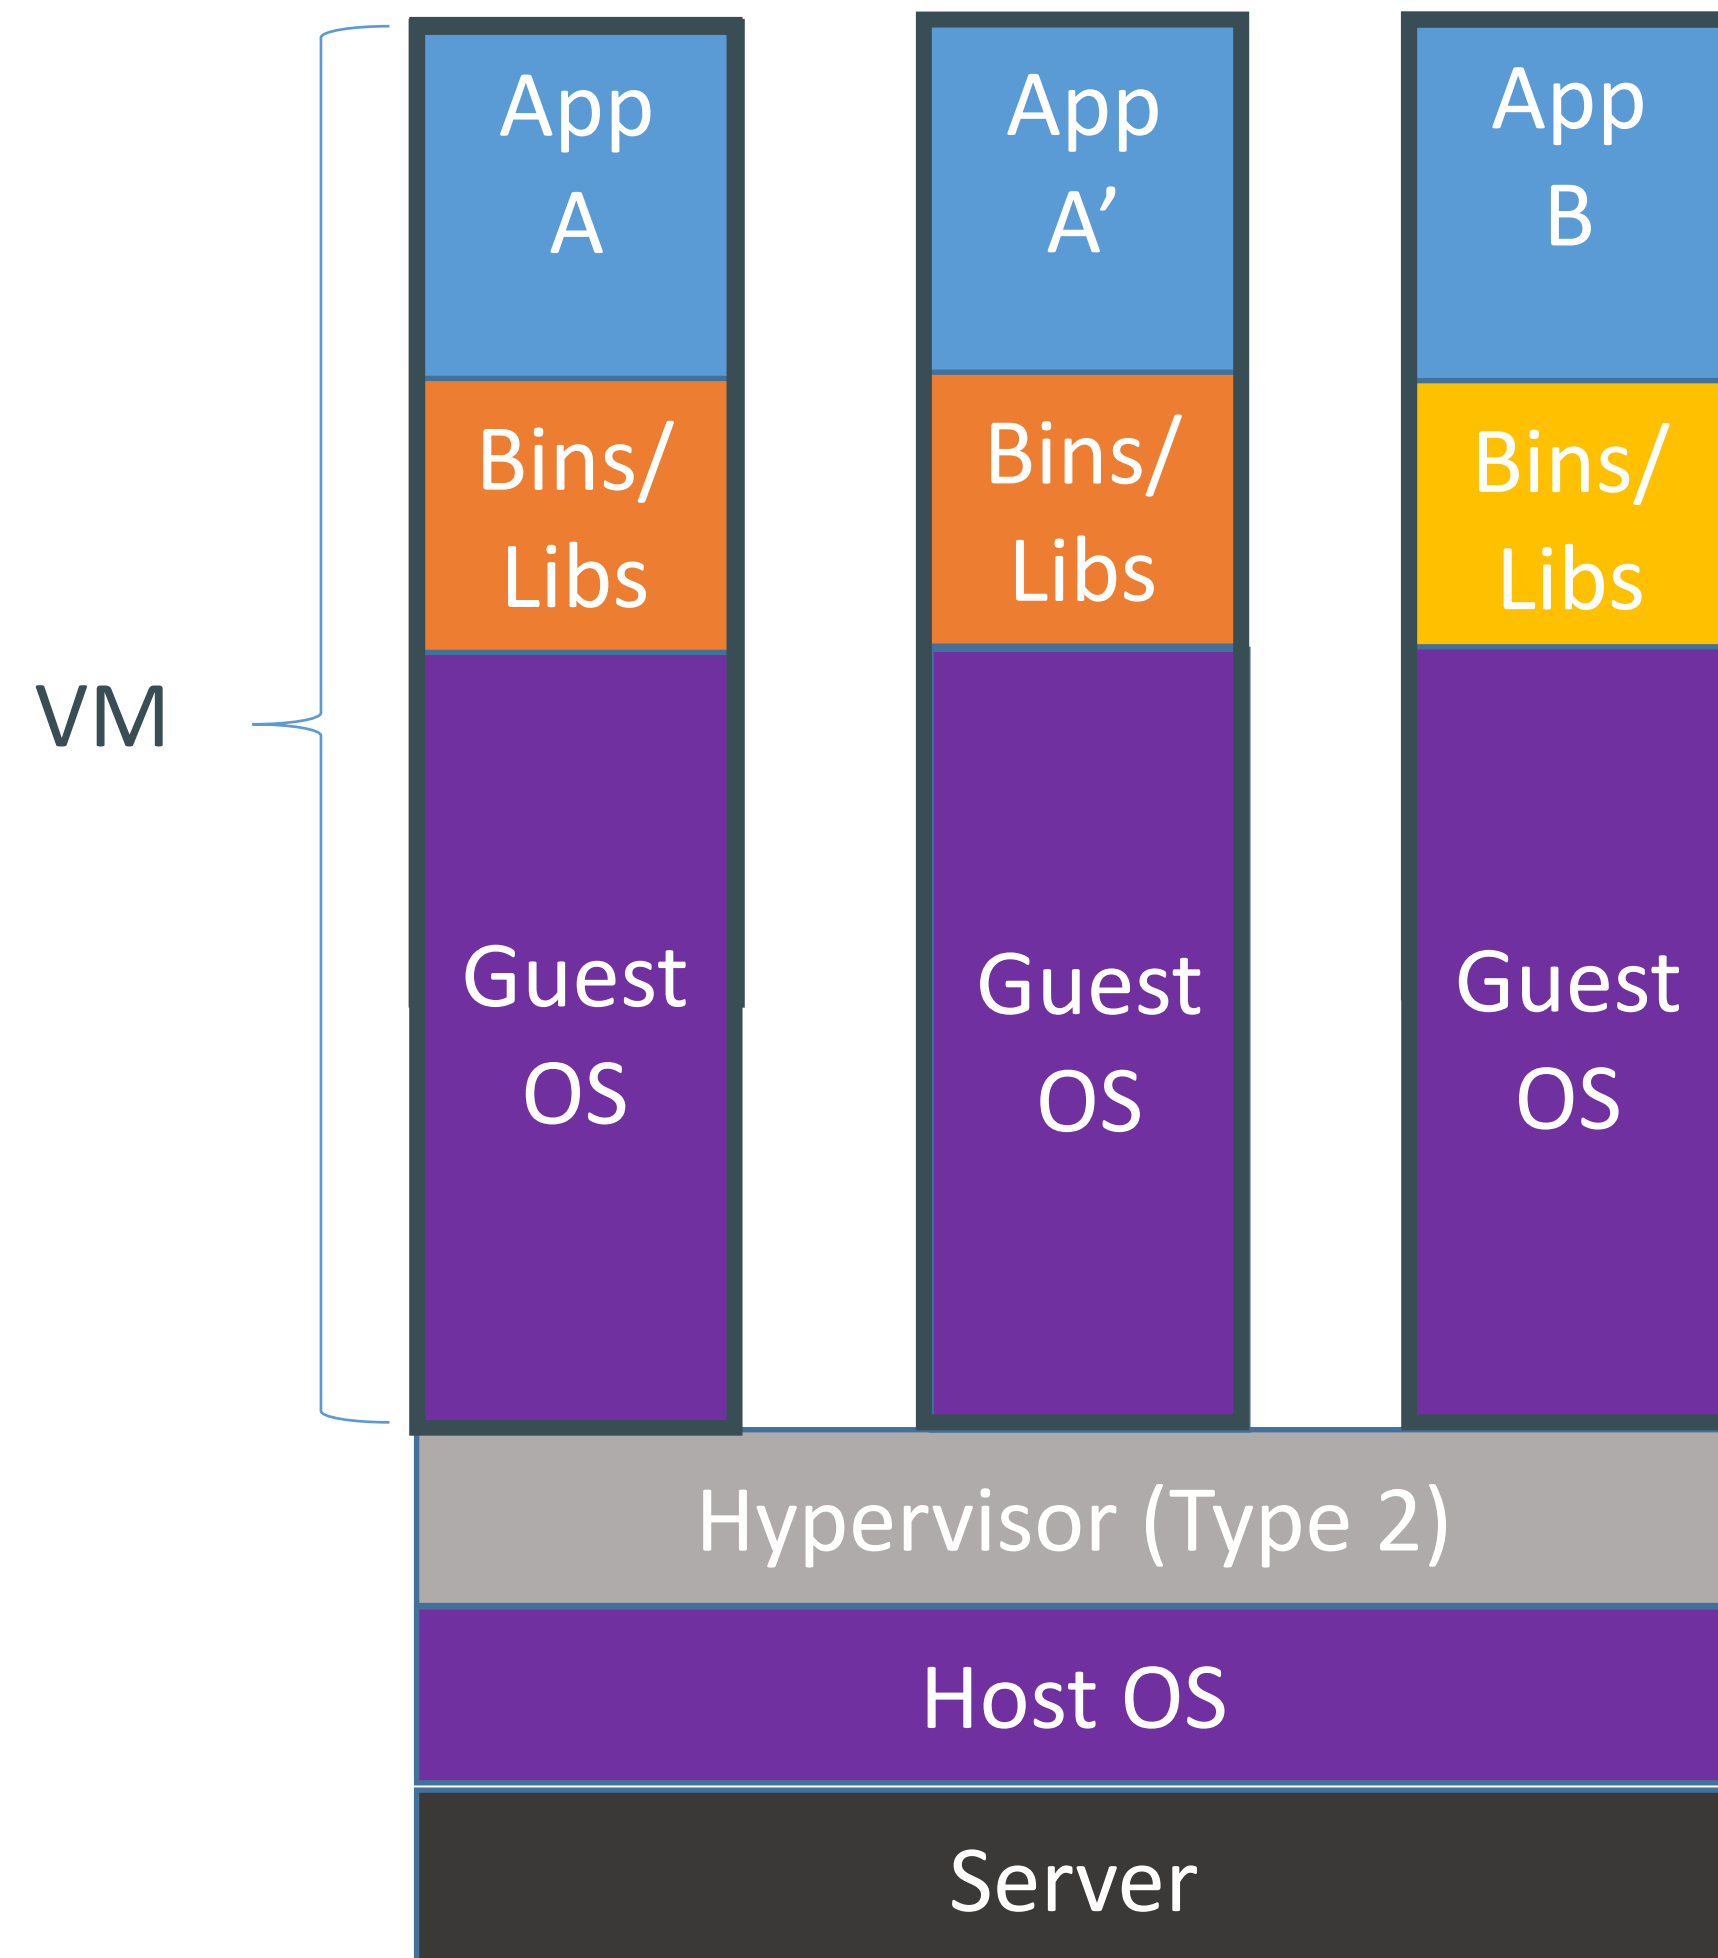
## WHY

- Run everywhere
  - Regardless of kernel version (2.6.32+)
  - Regardless of host distro
  - Physical or virtual, cloud or not
  - Container and host architecture must match\*
- Run anything
  - If it can run on the host, it can run in the container
  - i.e. if it can run on a Linux kernel, it can run

## WHAT

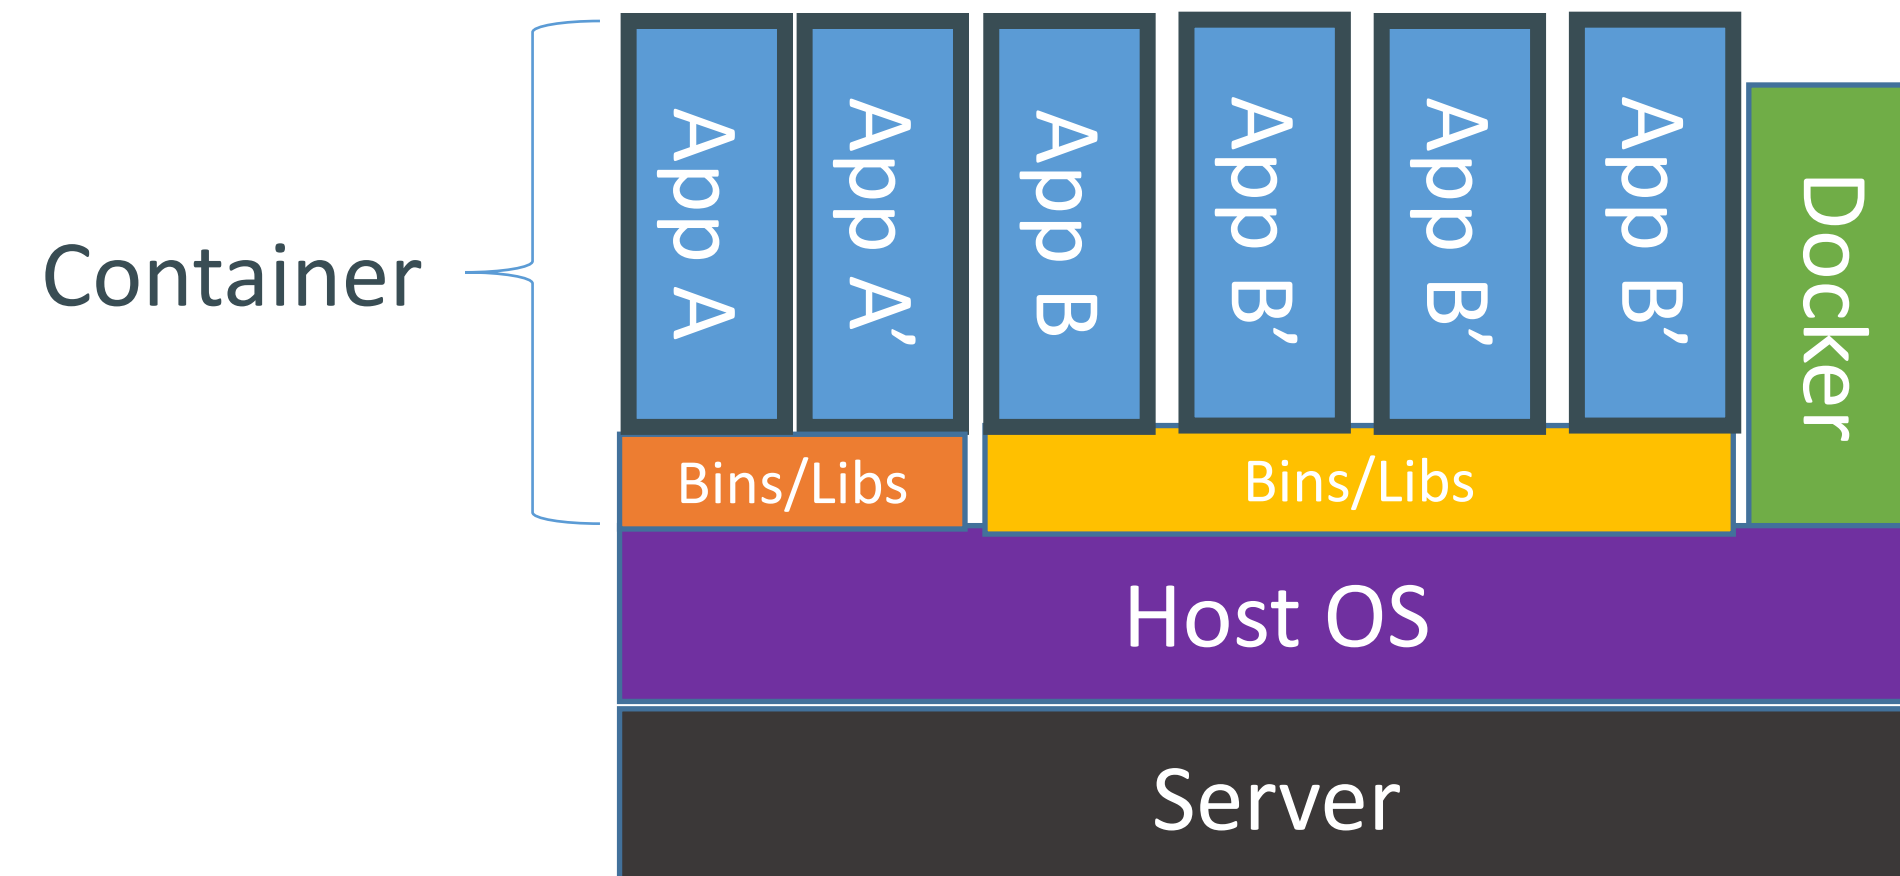
- High Level—It's a lightweight VM
  - Own process space
  - Own network interface
  - Can run stuff as root
  - Can have its own /sbin/init (different from host)
  - <<machine container>>
- Low Level—It's chroot on steroids
  - Can also *not* have its own /sbin/init
  - Container=isolated processes
  - Share kernel with host
  - No device emulation (neither HVM nor PV) from host)
  - <<application container>>

# Containers vs. VMs



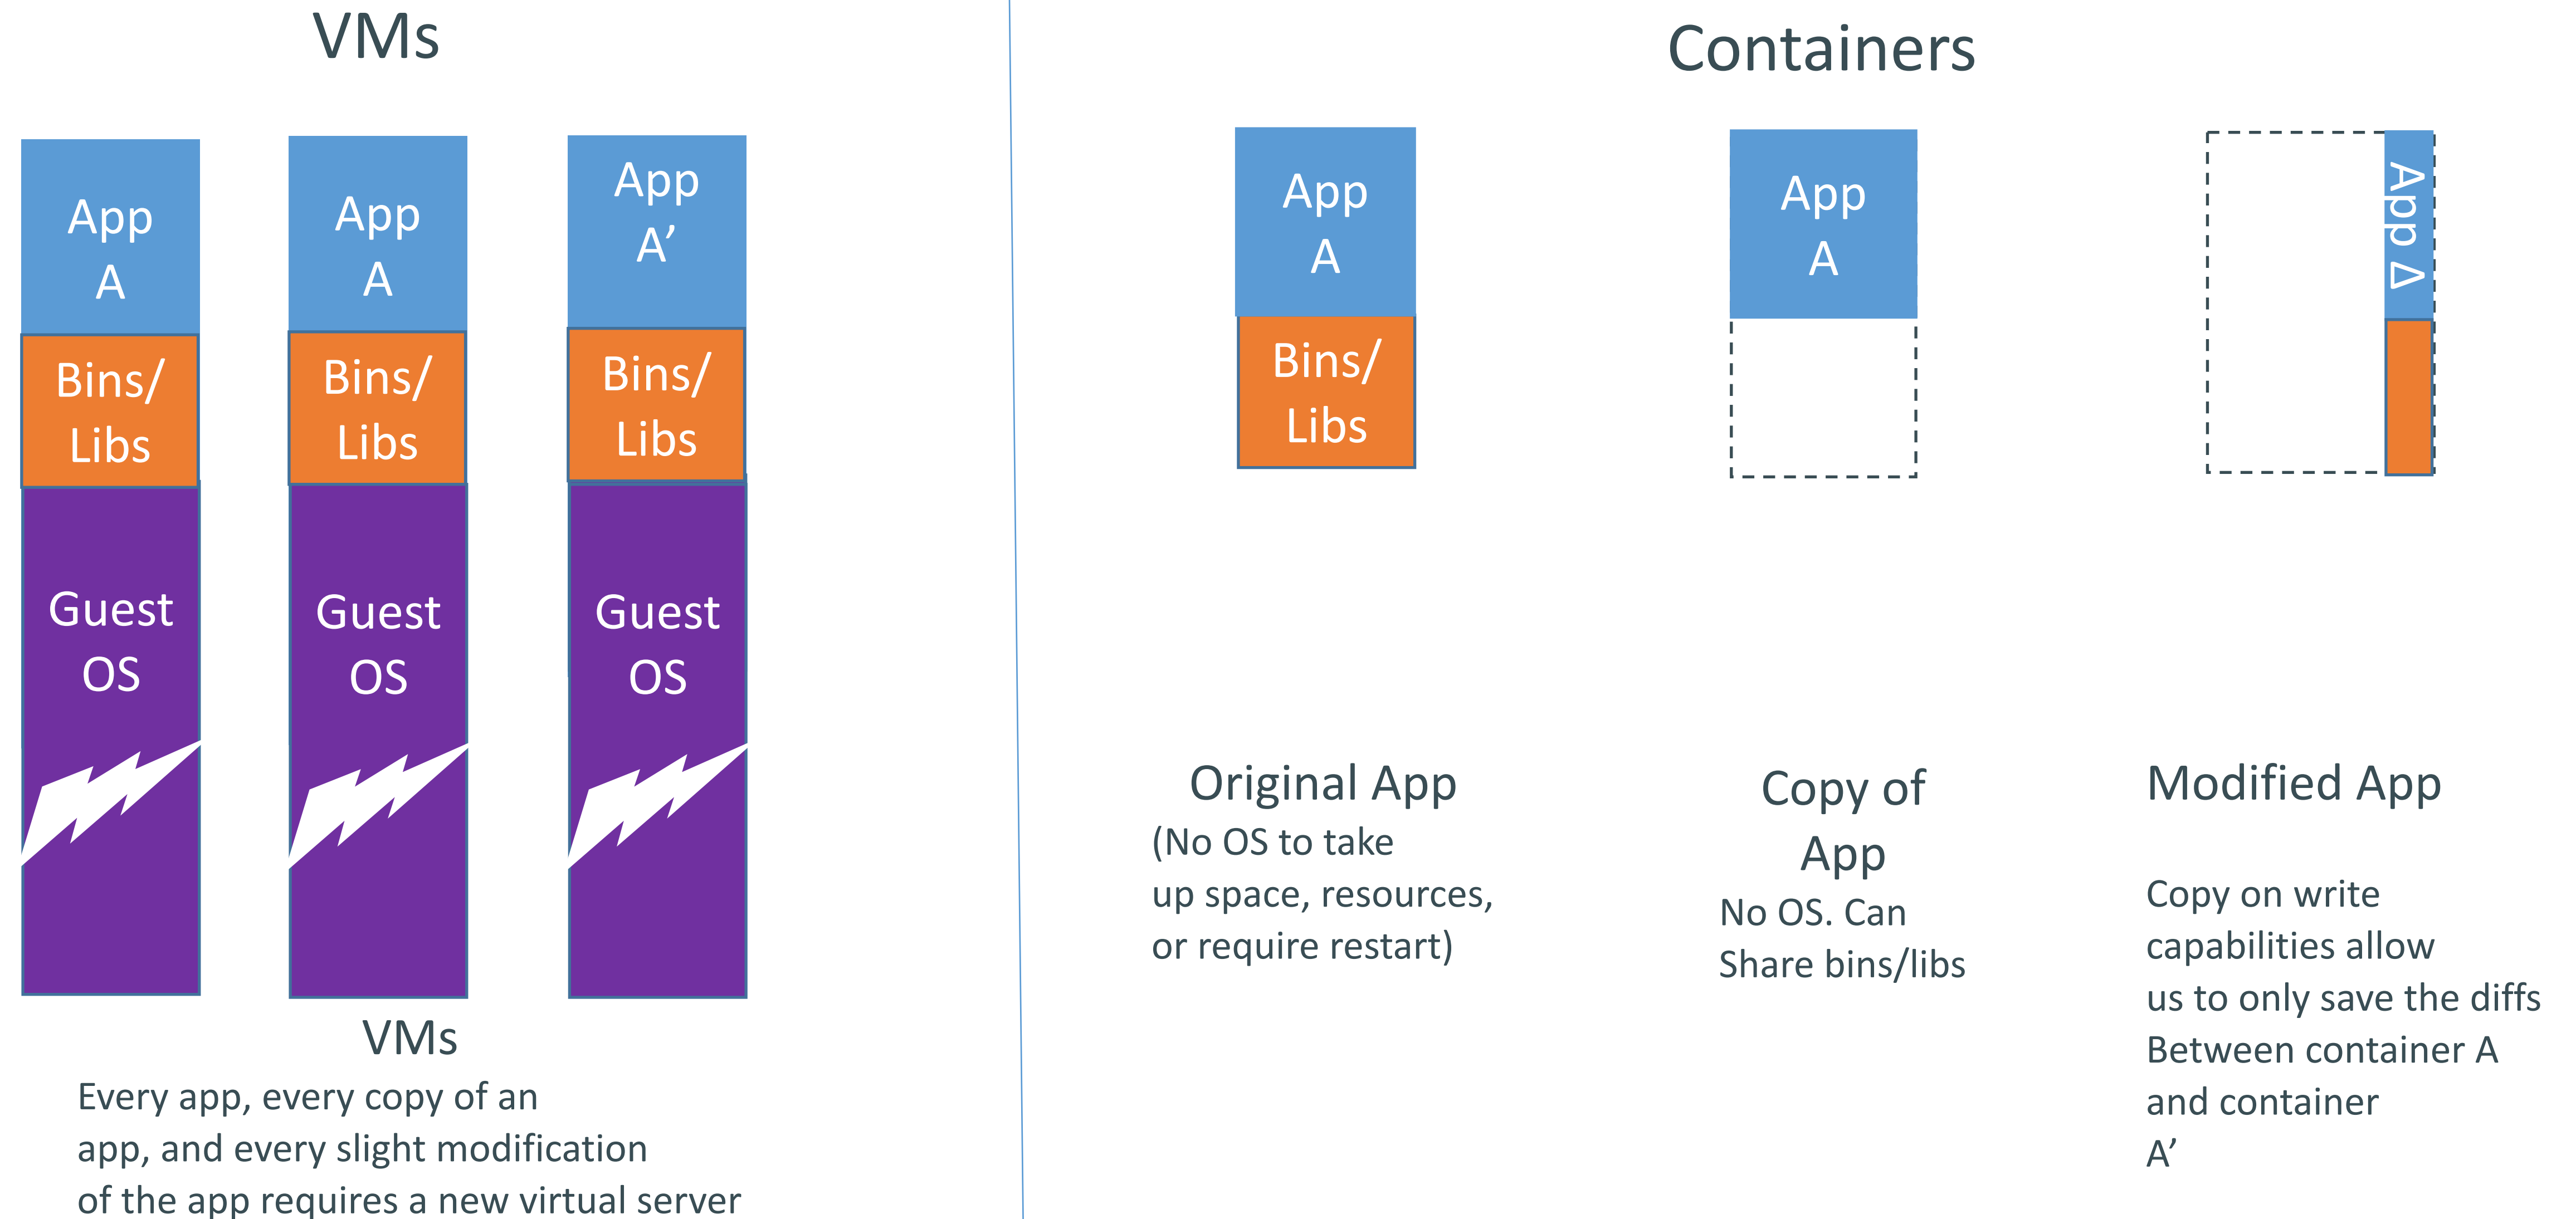
Containers are isolated, but share OS and, where appropriate, bins/libraries

...result is significantly faster deployment, much less overhead, easier migration, faster restart

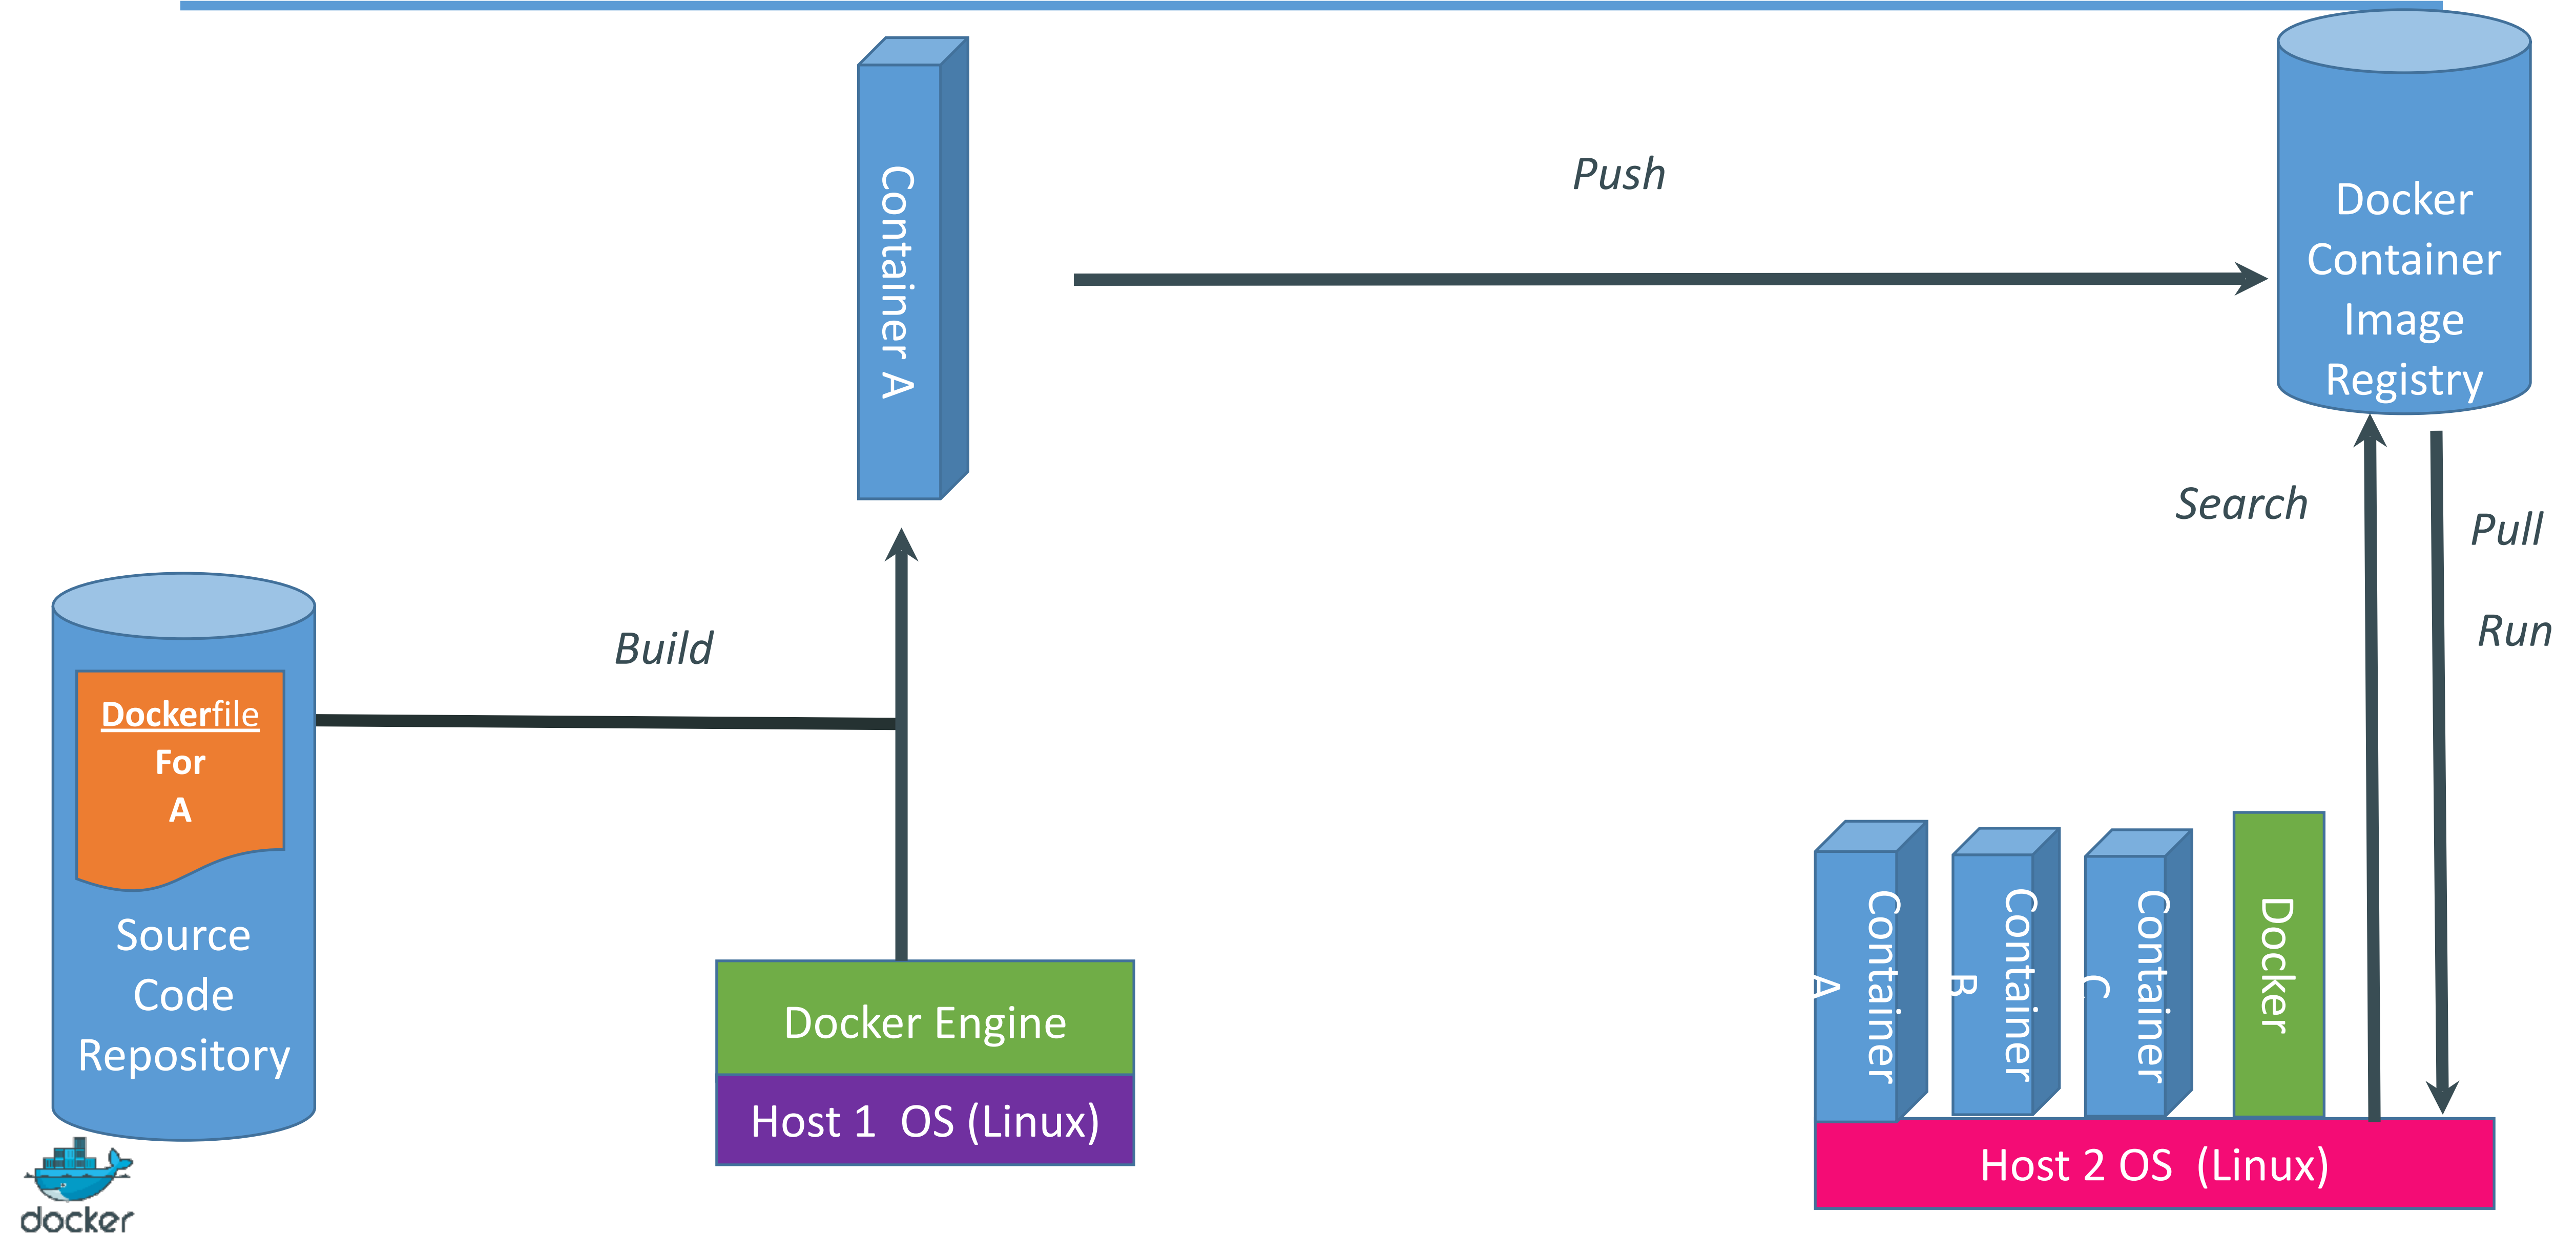




# Why are Docker containers lightweight?



# What are the basics of the Docker system?







# Ecosystem Support

---

- Operating systems
  - Virtually any distribution with a 2.6.32+ kernel
  - Red Hat/Docker collaboration to make work across RHEL 6.4+, Fedora, and other members of the family (2.6.32 +)
  - CoreOS—Small core OS purpose built with Docker
- OpenStack
  - Docker integration into NOVA (& compatibility with Glance, Horizon, etc.) accepted for Havana release
- Private PaaS
  - OpenShift
  - Solum (Rackspace, OpenStack)
  - Other TBA
- Public PaaS
  - Deis, Voxoz, Cocaine (Yandex), Baidu PaaS
- Public IaaS
  - Native support in Rackspace, Digital Ocean,+++
  - AMI (or equivalent) available for AWS & other
- DevOps Tools
  - Integrations with Chef, Puppet, Jenkins, Travis, Salt, Ansible +++
- Orchestration tools
  - Mesos, Heat, ++
  - Shipyard & others purpose built for Docker
- Applications
  - 1000's of Dockerized applications available at [index.docker.io](http://index.docker.io)



# Use Cases—From Docker Community

Use Case	Examples	Link
Clusters	Building a MongoDB cluster using docker	<a href="http://bit.ly/1acbjZf">http://bit.ly/1acbjZf</a>
	Production Quality MongoDB Setup with Docker	<a href="http://bit.ly/15CaiHb">http://bit.ly/15CaiHb</a>
	Wildfly cluster using Docker on Fedora	<a href="http://bit.ly/1bCIX0O">http://bit.ly/1bCIX0O</a>
Build your own PaaS	OpenSource PaaS built on Docker, Chef, and Heroku Buildpacks	<a href="http://deis.io">http://deis.io</a>
Web Based Environment for Instruction	JiffyLab – web based environment for the instruction, or lightweight use of, Python and UNIX shell	<a href="http://bit.ly/12oaj2K">http://bit.ly/12oaj2K</a>
Easy Application Deployment	Deploy Java Apps With Docker = Awesome	<a href="http://bit.ly/11BCvvu">http://bit.ly/11BCvvu</a>
	How to put your development environment on docker	<a href="http://bit.ly/1b4XtJ3">http://bit.ly/1b4XtJ3</a>
	Running Drupal on Docker	<a href="http://bit.ly/15MJS6B">http://bit.ly/15MJS6B</a>
	Installing Redis on Docker	<a href="http://bit.ly/16EWOKh">http://bit.ly/16EWOKh</a>
Create Secure Sandboxes	Docker makes creating secure sandboxes easier than ever	<a href="http://bit.ly/13mZGJH">http://bit.ly/13mZGJH</a>
Create your own SaaS	Memcached as a Service	<a href="http://bit.ly/11nL8vh">http://bit.ly/11nL8vh</a>
Automated Application Deployment	Multi-cloud Deployment with Docker	<a href="http://bit.ly/1bF3CN6">http://bit.ly/1bF3CN6</a>
Continuous Integration and Deployment	Next Generation Continuous Integration & Deployment with dotCloud's Docker and Strider	<a href="http://bit.ly/ZwTfoy">http://bit.ly/ZwTfoy</a>
	Testing Salt States Rapidly With Docker	<a href="http://bit.ly/1eFBtcm">http://bit.ly/1eFBtcm</a>
Lightweight Desktop Virtualization	<a href="http://bit.ly/14RYL6x">Docker Desktop: Your Desktop Over SSH Running Inside Of A Docker Container</a>	<a href="http://bit.ly/14RYL6x">http://bit.ly/14RYL6x</a>



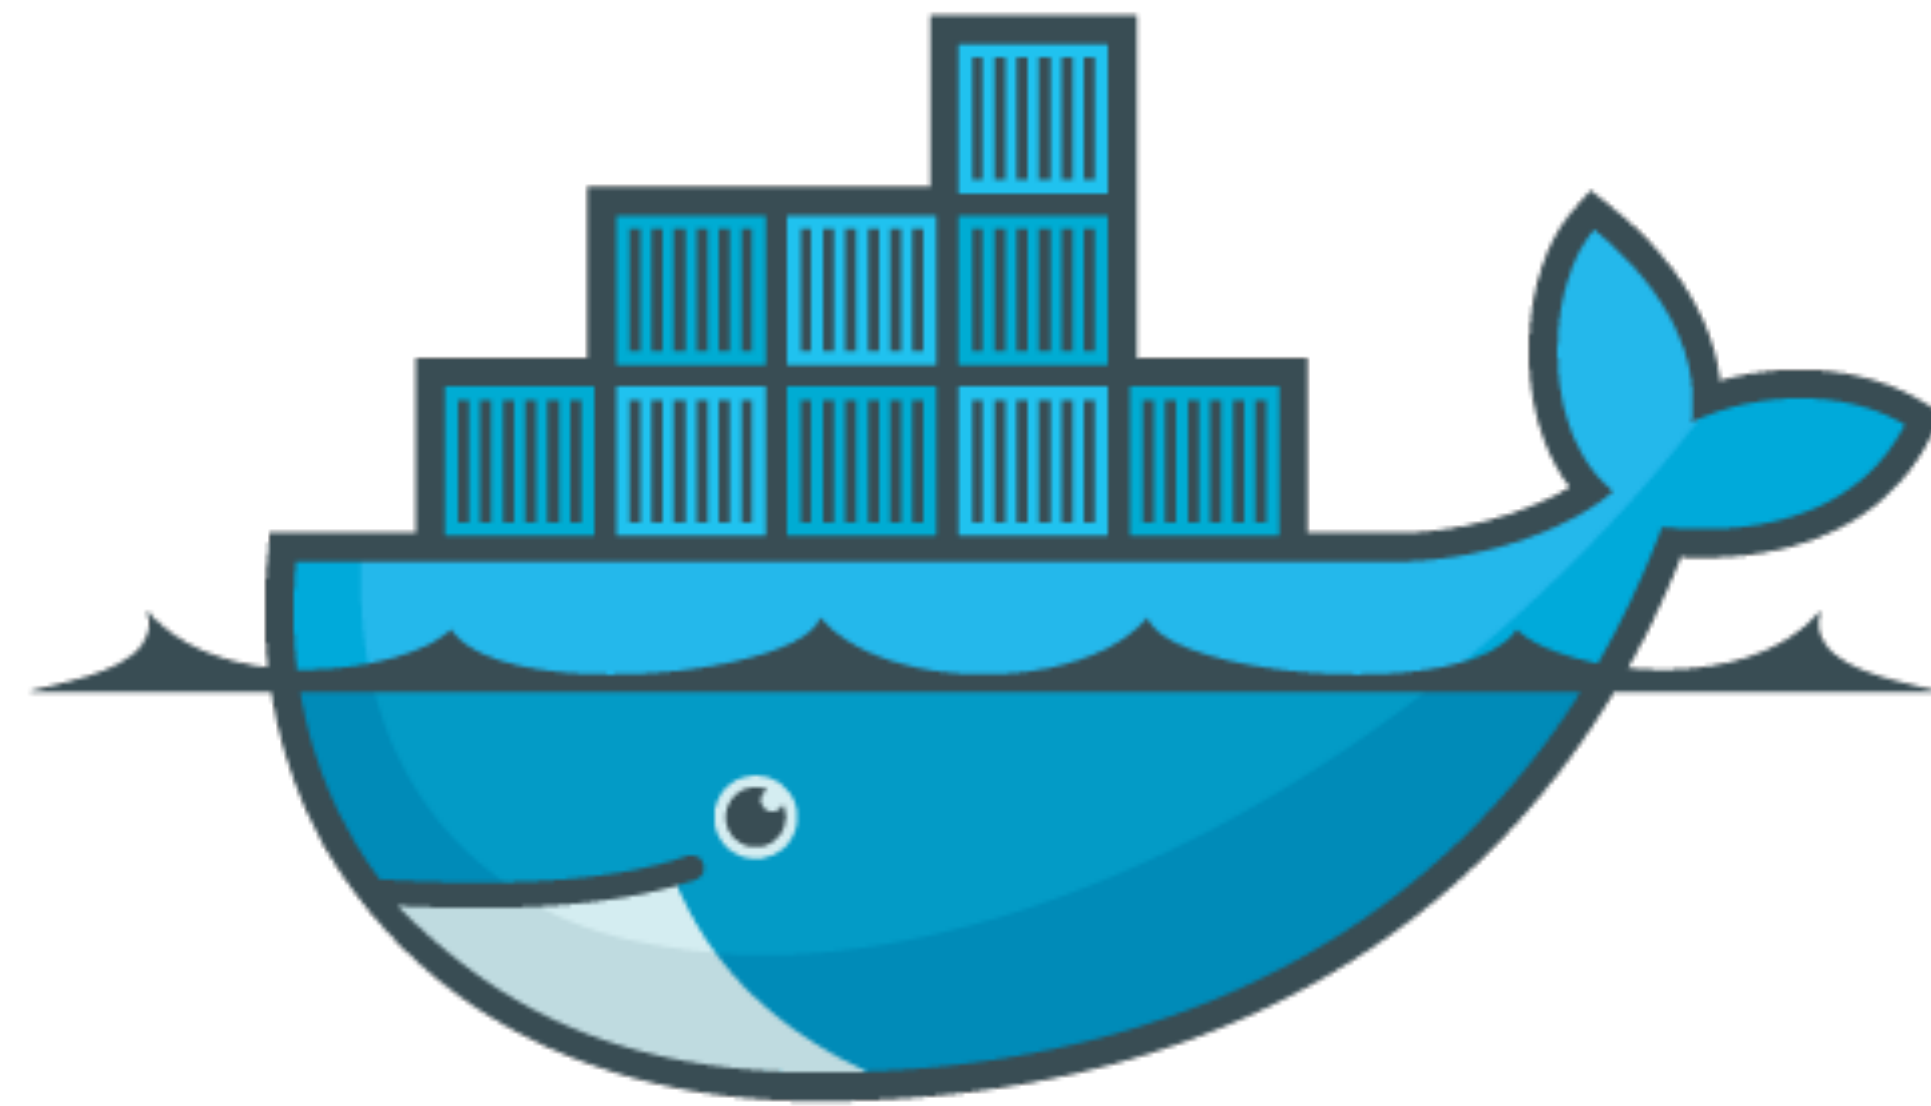
# Want to learn more?

---

- [www.docker.io](http://www.docker.io):
  - Documentation
  - Getting started: interactive tutorial, installation instructions, getting started guide,
  - About: Introductory whitepaper: <http://www.docker.io/the-whole-story/>
- Github: [dotcloud/docker](https://github.com/docker)
- IRC: [freenode/#docker](https://freenode.net/#docker)
- Google groups: [groups.google.com/forum/#!forum/docker-user](https://groups.google.com/forum/#!forum/docker-user)
- Twitter: follow @docker
- Meetups: Scheduled for Boston, San Francisco, Austin, London, Paris, Boulder...and Nairobi. <https://www.docker.io/meetups/>







docker  
[www.docker.io](http://www.docker.io)